

EDUCATION

University of Washington, Seattle

Sept. 2018—Est. Jun. 2022

B.S. in Computer Science

- Major GPA: 3.79 / 4.00
- Fields of Study: Programming Languages & Formal Verification & Compilers & MLSys
- Honors: Dean's List (College of Art & Science) 2018–Now; CRA Outstanding Undergraduate Researcher Award nominee of the Allen School (2021)

PUBLICATIONS

1. Marisa Kirisame*, Steven Lyubomirsky*, Altan Haan*, Jennifer Brennan, **Mike He**, Jared Roesch, Tianqi Chen, and Zachary Tatlock. Dynamic tensor rematerialization. In *International Conference on Learning Representations (ICLR'21)*, 2021 (*: Equal Contribution)
2. (**Under review at PLDI '22**) Bo-Yuan Huang, Steven Lyubomirsky, Yi Li, **Mike He**, Thierry Tambe, Gus Henry Smith, Akash Gaonkar, Vishal Canumalla, Gu-Yeon Wei, Aarti Gupta, Sharad Malik, and Zachary Tatlock. Specialized accelerators and compiler flows: Replacing accelerator apis with a formal software/hardware interface, 2021

WORKSHOP PAPERS

1. Bo-Yuan Huang*, Steven Lyubomirsky*, Thierry Tambe*, Yi Li, **Mike He**, Gus Smith, Gu-Yeon Wei, Aarti Gupta, Sharad Malik, and Zachary Tatlock. From dsls to accelerator-rich platform implementations: Addressing the mapping gap. In *Workshop on Languages, Tools, and Techniques for Accelerator Design (LATTE'21)*, 2021 (*: Equal Contribution)

EXPERIENCE

3LA, LATTE '21

June. 2020—Now

*Research Assistant @ PLSE**Seattle, WA*

- [3LA](#) proposes an end-to-end compilation flow that provides **flexible** and **verifiable** compiler support for custom Deep Learning (DL) accelerators. 3LA has a builtin implementation agnostic pattern matching algorithm that is capable of find accelerator supported workloads in DL models leveraging the power of Equality Saturation. Moreover, 3LA addresses the mapping gap between DL models represented in high-level domain-specific languages (DSLs) and specialized accelerators using Instruction-level Abstraction (**ILA**) as the software-hardware interface.
- **Talks & Presentations:**
 1. *From DSLs to Accelerator-rich Platform: Addressing the Mapping Gap*, Sept. 2021 at Intel (presented jointly with [Steven Lyubomirsky](#))
 2. *Correct & Flexible Compiler Support for Custom Accelerators*, Sept. 2021 at SRC ADA Center

Dynamic Tensor Rematerialization, ICLR '21

Oct. 2019—Aug. 2021

*Research Assistant @ PLSE**Seattle, WA*

- [Dynamic Tensor Rematerialization \(DTR\)](#) is a greedy gradient checkpointing algorithm. DTR **enables** training Deep Learning models on memory-constrained devices. Unlike previous approaches, DTR does not need any information of the DL model architectures ahead-of-time; instead it saves memory by evicting and recomputing tensors **on the fly**, i.e. trading time for memory, which further exploit opportunities of using gradient checkpointing on DL trainings. DTR is comparably efficient as previous approaches: it requires only $\mathcal{O}(N)$ more forward computations when training a N -layer linear feed-forward neural network with an $\Omega(\sqrt{N})$ memory budget.

Paul G. Allen School, University of Washington

Mar. 2021—June. 2021

*Teaching Assistant**Seattle, WA*

- Worked as a TA for **Principles of Programming Languages** (CSE 505)
- Helped re-designing CSE 505 and developing course materials for various topics about PL and formal verification (**Hoare Logic**, **Lambda Calculus** and **System F**, etc.) in **Coq**.

SERVICE

→ **MICRO '21**, Artifact Evaluation